

AI and Deep Learning in Finance

Sanjiv R. Das
Santa Clara University

ISB | FinTech | 2018

AI, ML and DL

Artificial Intelligence

Machine Learning

Deep Learning

The subset of machine learning composed of algorithms that permit software to train itself to perform tasks, like speech and image recognition, by exposing multilayered neural networks to vast amounts of data.

A subset of AI that includes abstruse statistical techniques that enable machines to improve at tasks with experience. The category includes deep learning

Any technique that enables computers to mimic human intelligence, using logic, if-then rules, decision trees, and machine learning (including deep learning)

Game Changers

- Mathematical innovations
- Hardware
- Big Data

BloombergMarkets ▾ Nasdaq's Latest Deal Shows That Data Reigns Supreme

Nasdaq's Latest Deal Shows That Data Reigns Supreme

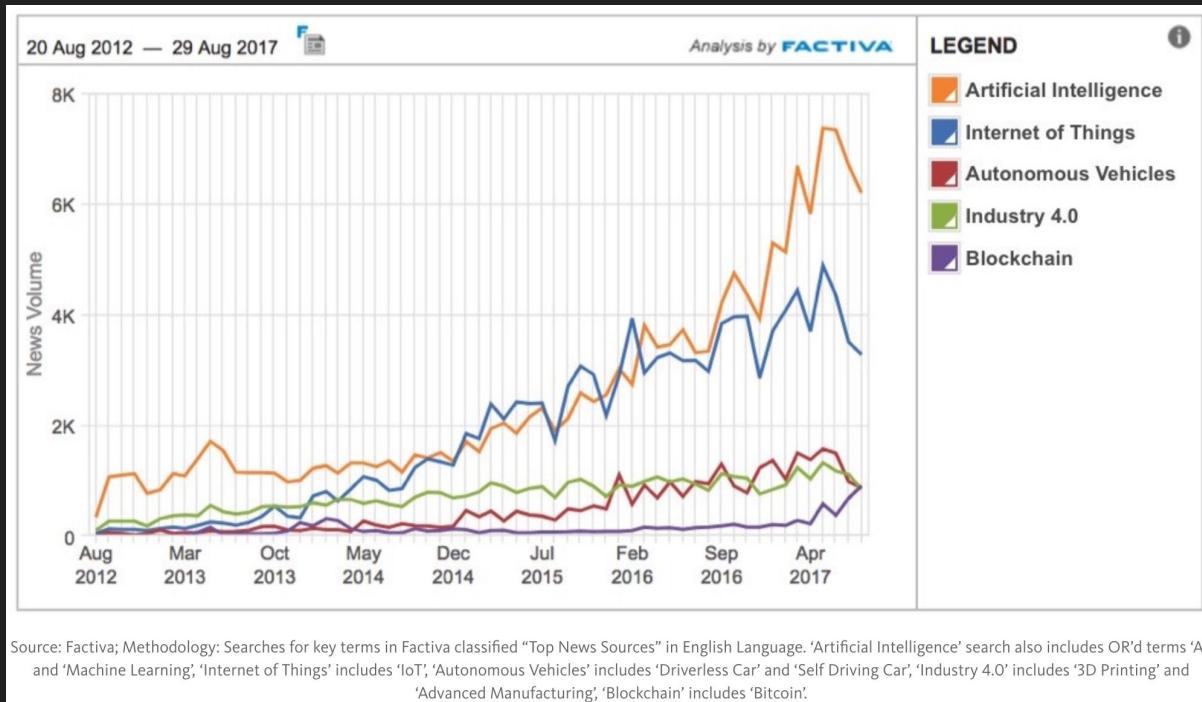
By **Annie Massa**
September 5, 2017, 7:33 AM PDT *Updated on* September 5, 2017, 12:57 PM PDT

→ Global exchange buys EVestment for \$705 million in debt, cash
→ EVestment offers specialized data for institutional investors

Nasdaq Inc.'s takeover of data provider EVestment Inc. spotlights the exchange industry's seemingly bottomless appetite for data and its emerging focus on cultivating a closer relationship with money managers.

AI and the Technological Singularity

<https://hackernoon.com/why-ai-is-now-on-the-menu-at-dinner-even-with-my-non-tech-friends-44c666348de4>



Transformation
vs
Change

Kurzweil claims that AI will pass the Turing test in 2029, and the singularity will come in 2045.

<https://futurism.com/kurzweil-claims-that-the-singularity-will-happen-by-2045/>

Huge Demand for AI Talent

<https://www.nytimes.com/2017/10/22/technology/artificial-intelligence-experts-salaries.html>

Home

The New York Times

TECHNOLOGY

Share

149

Tech Giants Are Paying Huge Salaries for Scarce A.I. Talent

Nearly all big tech companies have an artificial intelligence project, and they are willing to pay experts millions of dollars to help get it done.

[查看简体中文版](#) | [查看繁體中文版](#)

By CADE METZ OCT. 22, 2017



Definition of AI

- Intelligence exhibited by machines
- **Narrow or Weak AI**: “Expert systems that match or exceed human intelligence in a narrowly defined area, but not in broader areas” (Dvorsky G., 2013) e.g. Siri.
- **Artificial General Intelligence**: An artificial neural network not preprogrammed with fixed rules. Rewire itself to reflect patterns in the data, adaptable to its environment, in which (hopefully) advanced skills emerge organically.
- “Humans don’t learn to understand language by memorizing dictionaries and grammar books, so why should we possibly expect our computers to do so?” (LEWIS-KRAUS G, 2016).
- And, **Super AI**?

<http://io9.gizmodo.com/how-much-longer-before-our-first-ai-catastrophe-464043243> (Dvorsky G..
https://mobile.nytimes.com/2016/12/14/magazine/the-great-ai-awakening.html?_r=0&referer= (LEWIS-KRAUS G, 2016)

Two Types of AI

- Rule-Based AI
- Data-driven AI
- Example: Checkers. Albert Samuel @IBM began writing code for a checkers game program in 1949. In 1956, the program was demonstrated to the public on live television. In 1962, the computer beat checkers master player Robert Nealey, and IBM's stocks rose 15 percent overnight.
- Rule-based AI can never be more intelligent than its creators, but data-driven AI can!

Artificial Intelligence Learns to Learn Entirely on Its Own



A new version of AlphaGo needed no human instruction to figure out how to clobber the best Go player in the world — itself.

A mere 19 months after dethroning the world's top human Go player, the computer program AlphaGo has smashed an even more momentous barrier: It can now achieve unprecedented levels of mastery purely by teaching itself. Starting with zero knowledge of Go strategy and no training by humans, the new iteration of the program, called AlphaGo Zero, needed just three days to invent advanced strategies undiscovered by human players in the multi-millennia history of the game. By freeing artificial intelligence from a dependence on human knowledge, the breakthrough removes a primary limit on how smart machines can become.

Still it seems at least plausible to expect artificial intelligence in investing to outstrip human intelligence in discomfiting ways. If your computer figures out for itself how to invest using "advanced strategies undiscovered by human players," and if those strategies are so advanced that the computer can't even explain them to you, and if they make money, then what do you do? Do you trust that the computer knows what it's doing better than you do, and stand aside and give it all your money? Or do you worry that the computer's success reflects a particular market regime, or that the computer's strategy will break down once it starts working with a lot of money? Do you trust that the *computer* is thinking about these things, that it can see the big picture rather than just grabbing whatever works right now?

An Early History of AI

- Warren McCulloch and Walter Pitts (1943): proposed a model of artificial neurons in which each neuron is characterized as being "on" or "off" ; suggested that suitably defined networks could learn.
- The Turing Test (Alan Turing, 1950): designed to provide a satisfactory operational definition of intelligence. Intelligent behavior: the ability to achieve human-level performance in all cognitive tasks, sufficient to fool an interrogator.
- John McCarthy proposed the term ARTIFICIAL INTELLIGENCE (1956); defined the high-level language Lisp.
- Frank Rosenblatt (1962) - “perceptron convergence theorem”: his learning algorithm could adjust the connection strengths of a perceptron to match any input data.

An Early History of AI

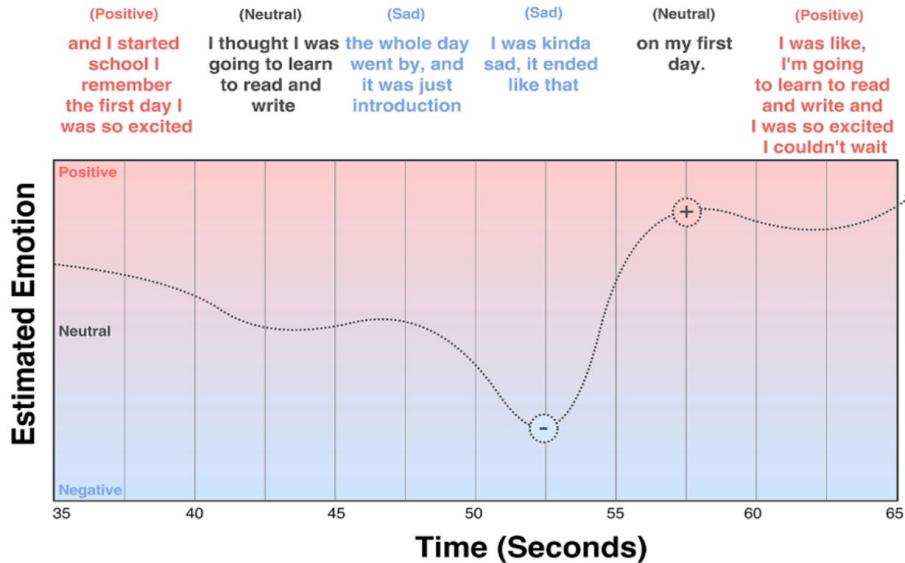
- Japanese "Fifth Generation" project (1981), a 10-year plan to build intelligent computers code; proposed to achieve full-scale natural language understanding.
- Fueled interest in the field in United States, led to the formation of The Microelectronics and Computer Technology Corporation (MCC) to counter the Japanese project.
- Reinvention of the back-propagation learning algorithm (1969) in mid 1980s.
- HITECH: the first computer program to defeat a grandmaster in a game of chess (Berliner, 1989).
- MARVEL: a real-time expert system that monitored massive data transmitted by a spacecraft, handling routine tasks and alerting the analysts to more serious problems (Schwuttke, 1992).
- PEGASUS: A speech understanding program to manage travel transactions (Zue et al., 1994).

Machine Learning

- Machine learning is the subfield of computer science that gives computers the ability to learn without being explicitly programmed." (Arthur Samuel)
- Definition (Tom Mitchell): "A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P if its performance at tasks in T , as measured by P , improves with experience E ."
- Supervised vs Unsupervised.

AI in Wearables

Emotion detectors Could Tell You When The Conversation Is Getting Awkward!



Graph showing real-time emotion detection [Graphic: courtesy of MIT CSAIL]

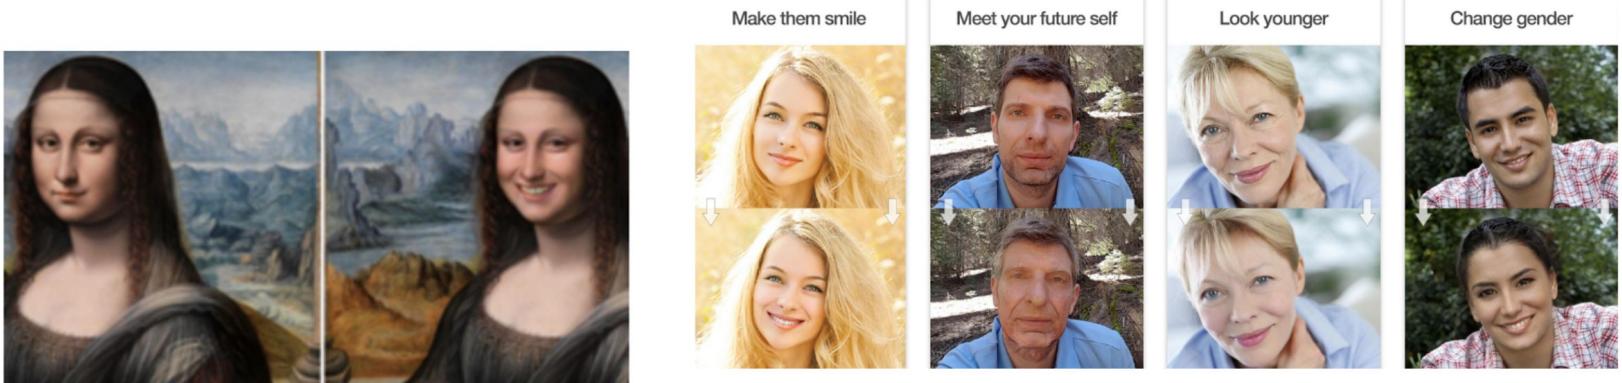


<https://www.fastcompany.com/3066867/wearables-that-detect-emotion-could-tell-you-when-the-conversation-is-getting-awkward>

FaceApp

FaceApp: AI Selfie Transformation App Can Even Make Mona Lisa Smile !

Uses neural networks to edit your selfie via photo-realistic filters



https://news.developer.nvidia.com/this-ai-selfie-transformation-app-can-even-make-mona-lisa-smile/?_lrcsc=a06a736b-bf67-4a30-b40f-368d4225b35b&ncid=so-lin-It-798

Conversational AI

TalkIQ

CH⁺RUS


TROOPS

NarrativeScience 



P

 conversica

Chat Bots



AMAZON'S ALEXA



GOOGLE'S ASSISTANT



APPLE'S SIRI

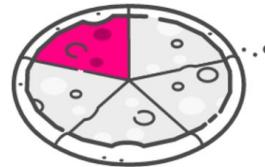


MICROSOFT'S CORTANA

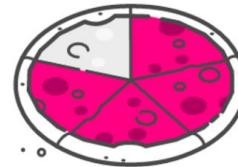
Insurance

Lemonade

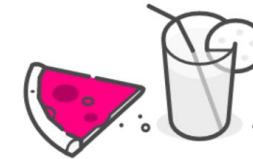
- Improving fairness of the insurance industry through Maya, their artificial intelligence bot
- Social Impact is the legal mission and business model
- Faster, easier



A transparent 20%
fee to run everything



We pay claims
super fast



If there's money leftover,
we give it back to causes

<https://www.lemonade.com/>

Healthcare

Remedy

- Improving primary health care treatment
- Remy: The world's smartest medical AI
- Taught by physicians, driven by data
- No insurance required



Deep Patient:

- **An Unsupervised Representation**
- **to Predict the Future of Patients**
- **from the Electronic Health Records**
- discovered patterns hidden in the hospital data to indicate
- people's disposition to a wide range of ailments,
- including cancer of the liver

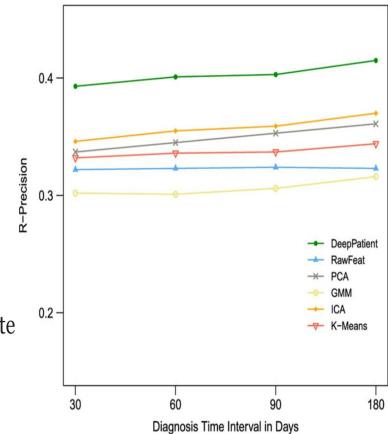


Figure 3: R-precision obtained in the disease tagging experiment by the different patient representations over several prediction time intervals (expressed as number of days).

<https://www.remedymedical.com/>

<https://www.nature.com/articles/srep26094>

Estimating the effects of technology

- (Roy) Amara's Law: "We tend to overestimate the effect of a technology in the short run and underestimate the effect in the long run."
- Arthur C. Clarke's Three Laws:
 - a. When a distinguished but elderly scientist states that something is possible, he is almost certainly right. When he states that something is impossible, he is very probably wrong.
 - b. The only way of discovering the limits of the possible is to venture a little way past them into the impossible.
 - c. Any sufficiently advanced technology is indistinguishable from magic.

Deep Learning in Finance

- Bridgewater Associates: World's largest hedge fund has a project to automate decision-making to save time and eliminate human emotion volatility.
- Goldman Sachs: Two out of the 600 equity traders left. Found that four traders can be replaced by one computer engineer.

Deep Learning in Finance

- Transactions: By 2020 at least five percent of all economic transactions will be handled by autonomous software. AI will process payment functions and learn from customer behaviours, through Intelligent Payment Management (IPM).
- Savings: AI will help consumers make daily financial decisions and monitor spending. New Personal Financial Management apps use contextual awareness, which measures spending habits and online footprints to create personalised advice. Combining pooled financial data with end-user control to offer tailor-made services is a classic AI solution.

Deep Learning in Finance

- Cross-selling: Categorization-as-a-Service (Caas), for understanding customer transactions for cross selling.
- Fraud prevention: Mining user data to detect abnormal behavior, anomalies, and unusual transactions.

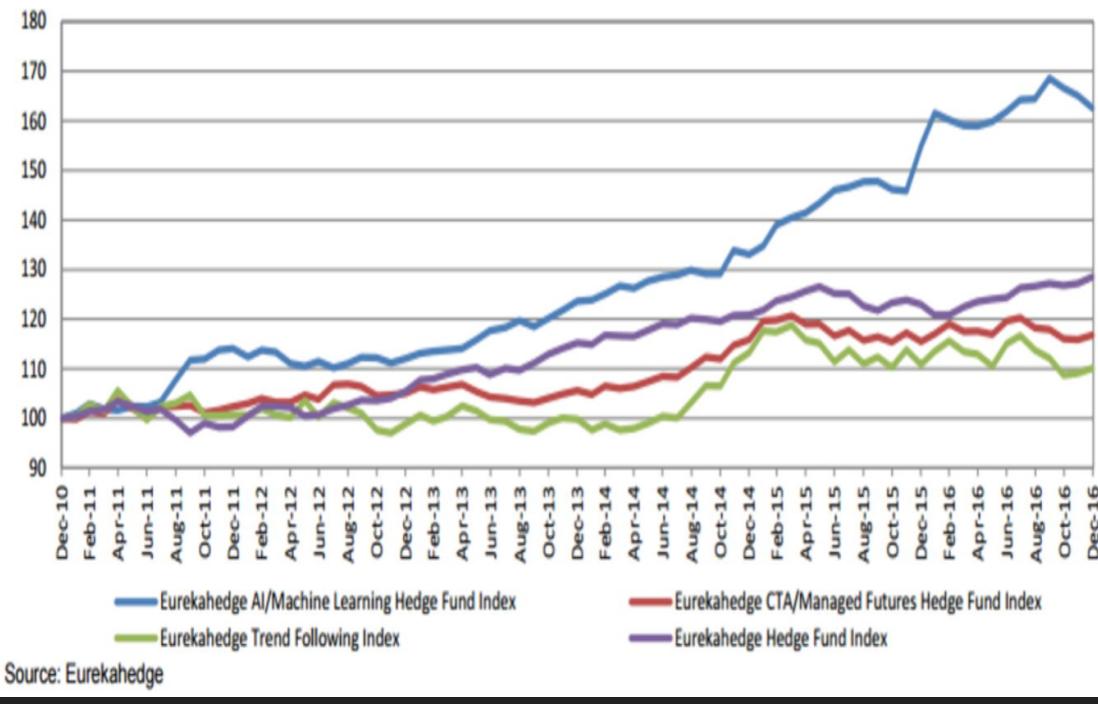
Deep Learning in Finance

- Mizuho Financial Group sent Pepper, its humanoid robot into its Tokyo branch to handle customer inquiries. Partnering with IBM to enable Pepper to understand human emotions, and build interaction into apps.
- RBS is trialing Luvo AI, a customer service assistant to interact with staff and customers.

Deep Learning in Finance

- AXA (insurer) has an app-based bot called Xtra, it engages in bespoke conversations with customers about healthy living.
- AI is used in peer2peer lending.

Hedge Funds use Machine Learning



Blackrock: replacing human stock pickers with machine algorithms.

Sentient Inc: Hedge fund run entirely using AI. Secret algo with adaptive learning. Uses thousands of machines.

Numerai: Hedge fund makes trades by aggregating trading algorithms submitted by anonymous contributors, prizes are awarded in cryptocurrency.

Very little data about the track record of these hedge funds, as the business remains secretive.

Investor reluctance to turn over money completely to a machine.

Emma: Evolved a hedge fund using a software that writes news articles.

Financial Automation

BloombergMarkets ▾

JPMorgan Software Does in Seconds What Took Lawyers 360,000 Hours

JPMorgan Software Does in Seconds What Took Lawyers 360,000 Hours

by **Hugh Son**

February 27, 2017, 4:31 PM PST Updated on February 28, 2017, 4:24 AM PST

At JPMorgan Chase & Co., a learning machine is parsing financial deals that once kept legal teams busy for thousands of hours.

The program, called COIN, for Contract Intelligence, does the mind-numbing job of interpreting commercial-loan agreements that, until the project went online in June, consumed 360,000 hours of work each year by lawyers and loan officers. The software reviews documents in seconds, is less error-prone and never asks for vacation.

Deep Learning is Pattern Recognition

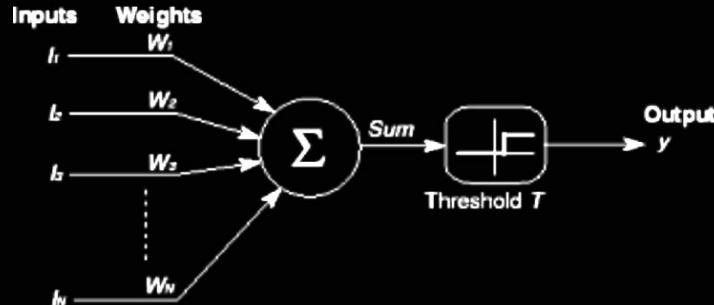


Figure 1.1: McCullough Pitts neuron

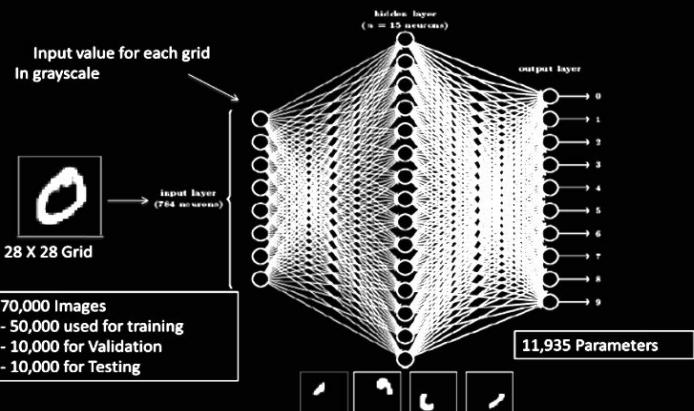


Figure 2.2: The NN used for Solving the MNIST Digit Recognition Problem

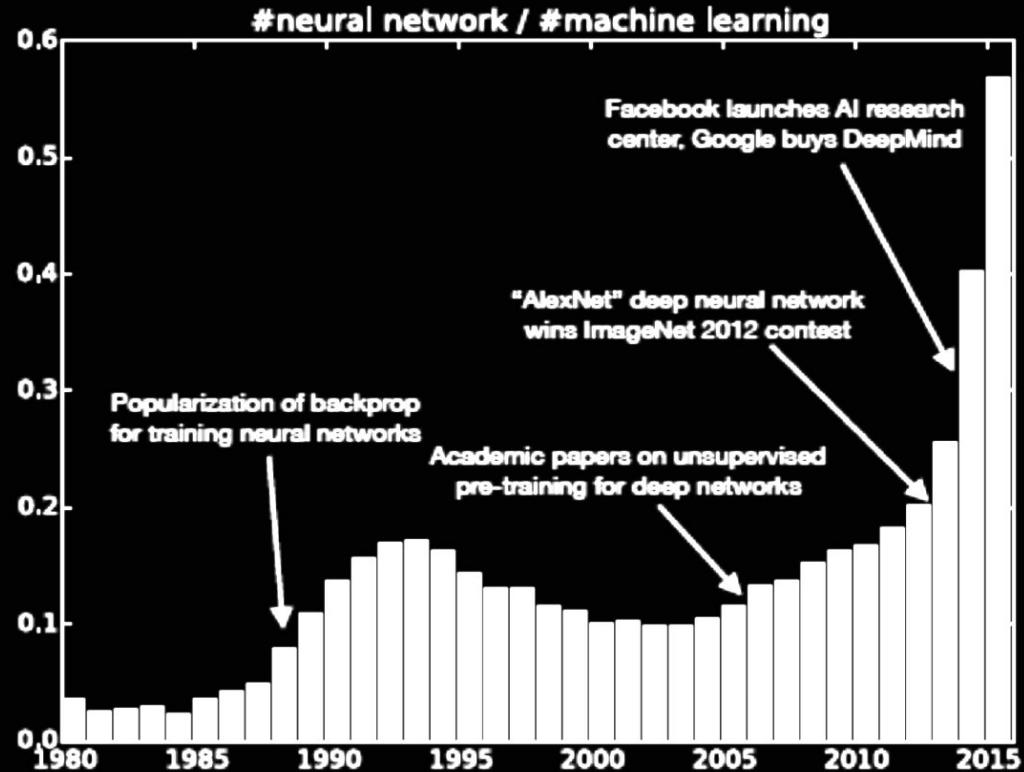
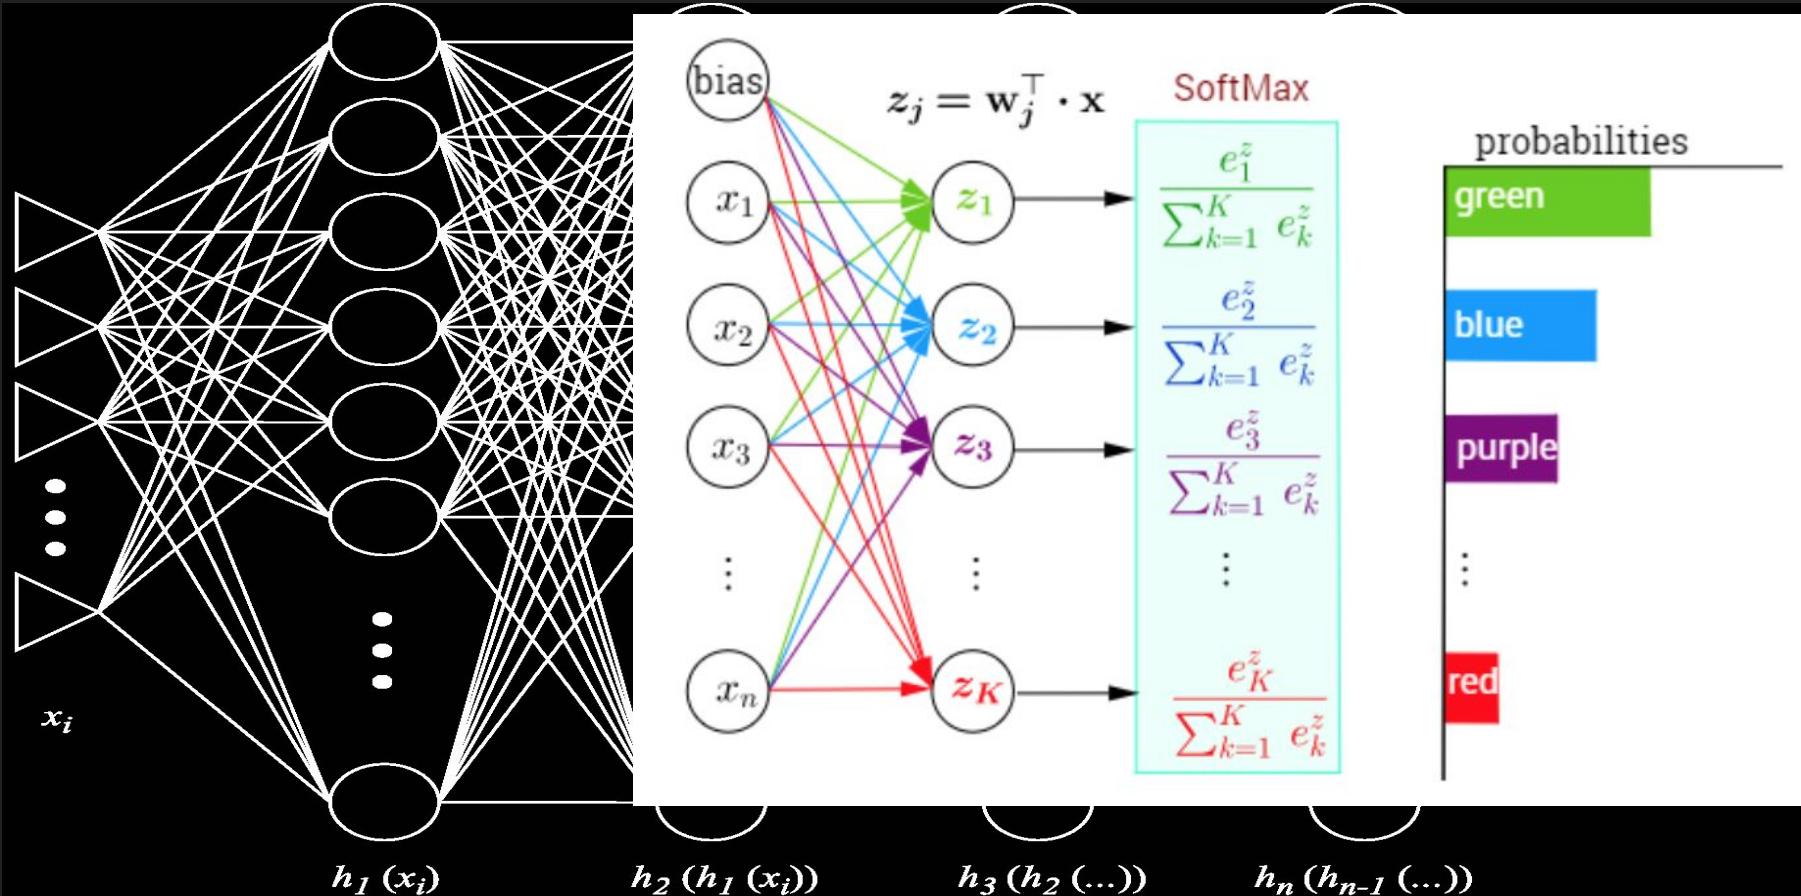
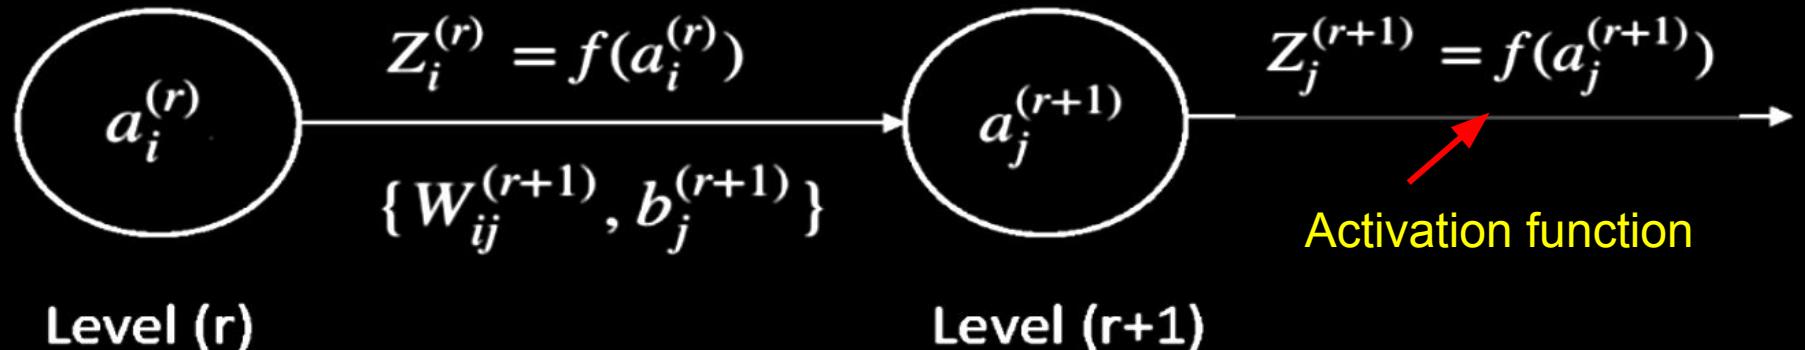


Figure 1.2: Number of NN research reports

The Mathematics of Deep Learning (<http://srdas.github.io/DLBook/>)



Subset of the Net



Level (r)

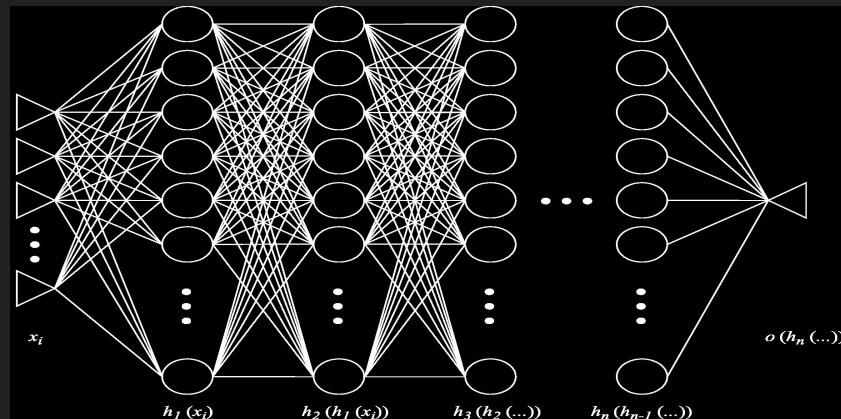
Level (r+1)

$$a_j = b_j + \sum_{I=1}^n w_{ij} z_i$$

Net input

$$\sigma(a_j) = \frac{1}{1 + \exp(-a_j)}$$

Sigmoid function



Activation Functions

https://en.wikipedia.org/wiki/Activation_function

Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer NN	

Loss Function

Fitting the DLN is an exercise where the best weights $\{W, b\} = \{W_{ij}^{(r+1)}, b_j^{(r+1)}\}$, $\forall r$ for all layers are determined to minimize a loss function generally denoted as

$$\min_{W,b} \sum_{m=1}^M L_m[h(X^{(m)}), T^{(m)}]$$

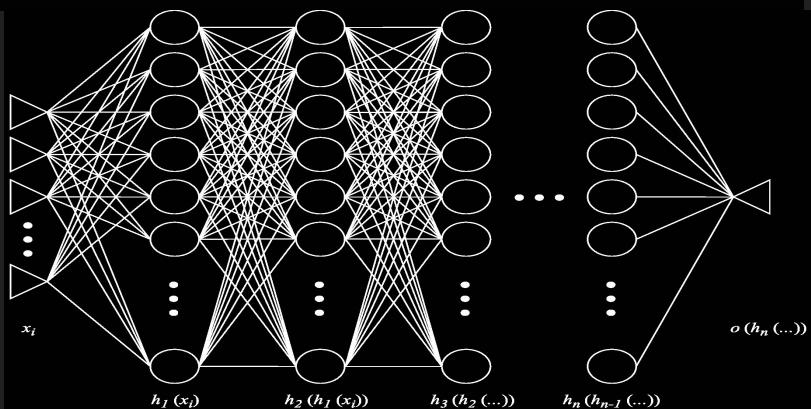
where M is the number of training observations (rows in the data set), $T^{(m)}$ is the true value of the output, and $h(X^{(m)})$ is the model output from the DLN. The loss function L_m quantifies the difference between the model output and the true output.

Cross
Entropy

$$C = -\frac{1}{n} \sum_x [y \ln a + (1 - y) \ln(1 - a)]$$

Quadratic
Loss

$$C = \frac{(y - a)^2}{2}$$



Gradient Descent

Fitting the DLN requires getting the weights $\{W, b\}$ that minimize L_m . These are done using gradient descent, i.e.,

$$W_{ij}^{(r+1)} \leftarrow W_{ij}^{(r+1)} - \eta \cdot \frac{\partial L_m}{\partial W_{ij}^{(r+1)}}$$

$$b_j^{(r+1)} \leftarrow b_j^{(r+1)} - \eta \cdot \frac{\partial L_m}{\partial b_j^{(r+1)}}$$

Here η is the learning rate parameter. We iterate on these functions until the gradients become zero, and the weights discontinue changing with each update, also known as an **epoch**.

Total Gradient

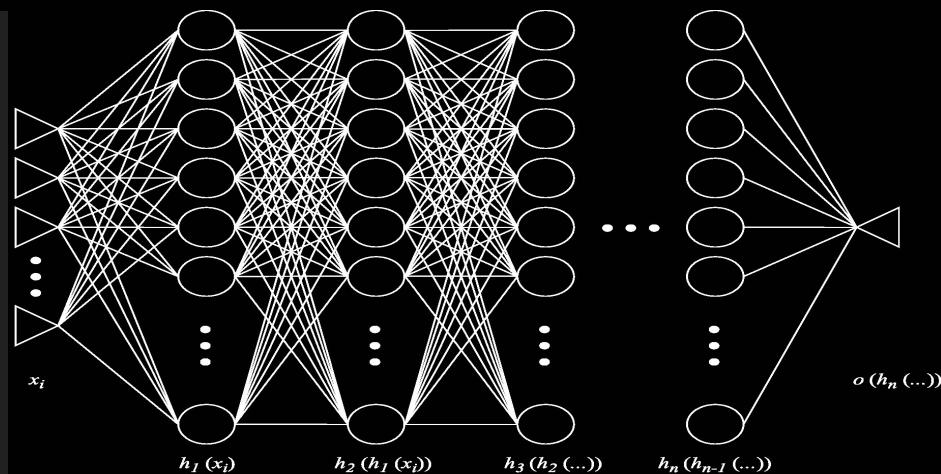
$$w_{ij}^{(r)} \leftarrow w_{ij}^{(r)} - \frac{\eta}{M} \sum_{m=1}^M \frac{\partial L_m(w, b; X(m), T(m))}{\partial w_{ij}^{(r)}}$$

$$b_i^{(r)} \leftarrow b_i^{(r)} - \frac{\eta}{M} \sum_{m=1}^M \frac{\partial L_m(w, b; X(m), T(m))}{\partial b_i^{(r)}}$$

Stochastic Batch Gradient

$$w_{ij}^{(r)} \leftarrow w_{ij}^{(r)} - \frac{\eta}{s} \sum_{m=1}^s \frac{\partial L_m(w, b; X(m), T(m))}{\partial w_{ij}^{(r)}}$$

$$b_i^{(r)} \leftarrow b_i^{(r)} - \frac{\eta}{s} \sum_{m=1}^s \frac{\partial L_m(w, b; X(m), T(m))}{\partial b_i^{(r)}}$$



Gradients and the Chain Rule

To solve this minimization problem, we need gradients for all W, b . These are denoted:

$$a_j^{(r+1)} = \sum_{i=1}^{n_r} W_{ij}^{(r+1)} Z_i^{(r)} + b_j^{(r+1)}$$

$$\frac{\partial L_m}{\partial W_{ij}^{(r+1)}},$$

$$\frac{\partial L_m}{\partial b_j^{(r+1)}},$$

$\forall r+1, j$

We write out these gradients using the chain rule:

$$\frac{\partial L_m}{\partial W_{ij}^{(r+1)}} = \frac{\partial L_m}{\partial a_j^{(r+1)}} \cdot \frac{\partial a_j^{(r+1)}}{\partial W_{ij}^{(r+1)}} = \delta_j^{(r+1)} \cdot Z_i^{(r)}$$

where we have written

$$\delta_j^{(r+1)} = \frac{\partial L_m}{\partial a_j^{(r+1)}}$$

Likewise, we have

$$\frac{\partial L_m}{\partial b_j^{(r+1)}} = \frac{\partial L_m}{\partial a_j^{(r+1)}} \cdot \frac{\partial a_j^{(r+1)}}{\partial b_j^{(r+1)}} = \delta_j^{(r+1)} \cdot 1 = \delta_j^{(r+1)}$$

Delta Values

we need the following intermediate calculation

$$\begin{aligned} a_j^{(r+1)} &= \sum_{i=1}^{n_r} W_{ij}^{(r+1)} Z_i^{(r)} + b_j^{(r+1)} \\ &= \sum_{i=1}^{n_r} W_{ij}^{(r+1)} f(a_i^{(r)}) + b_j^{(r+1)} \end{aligned}$$

$$\frac{\partial a_j^{(r+1)}}{\partial a_i^{(r)}} = W_{ij}^{(r+1)} \cdot f'(a_i^{(r)})$$

$$\delta_i^{(r)} = \frac{\partial L_m}{\partial a_i^{(r)}}$$

$$= \sum_{j=1}^{n_{r+1}} \frac{\partial L_m}{\partial a_j^{(r+1)}} \cdot \frac{\partial a_j^{(r+1)}}{\partial a_i^{(r)}}$$

$$= \sum_{j=1}^{n_{r+1}} \delta_j^{(r+1)} \cdot W_{ij}^{(r+1)} \cdot f'(a_i^{(r)})$$

$$= f'(a_i^{(r)}) \cdot \sum_{j=1}^{n_{r+1}} \delta_j^{(r+1)} \cdot W_{ij}^{(r+1)}$$

Output Layer

$$a_j^{(R+1)} = \sum_{i=1}^{n_R} W_{ij}^{(R+1)} Z_j^{(R)} + b_j^{(R+1)}$$

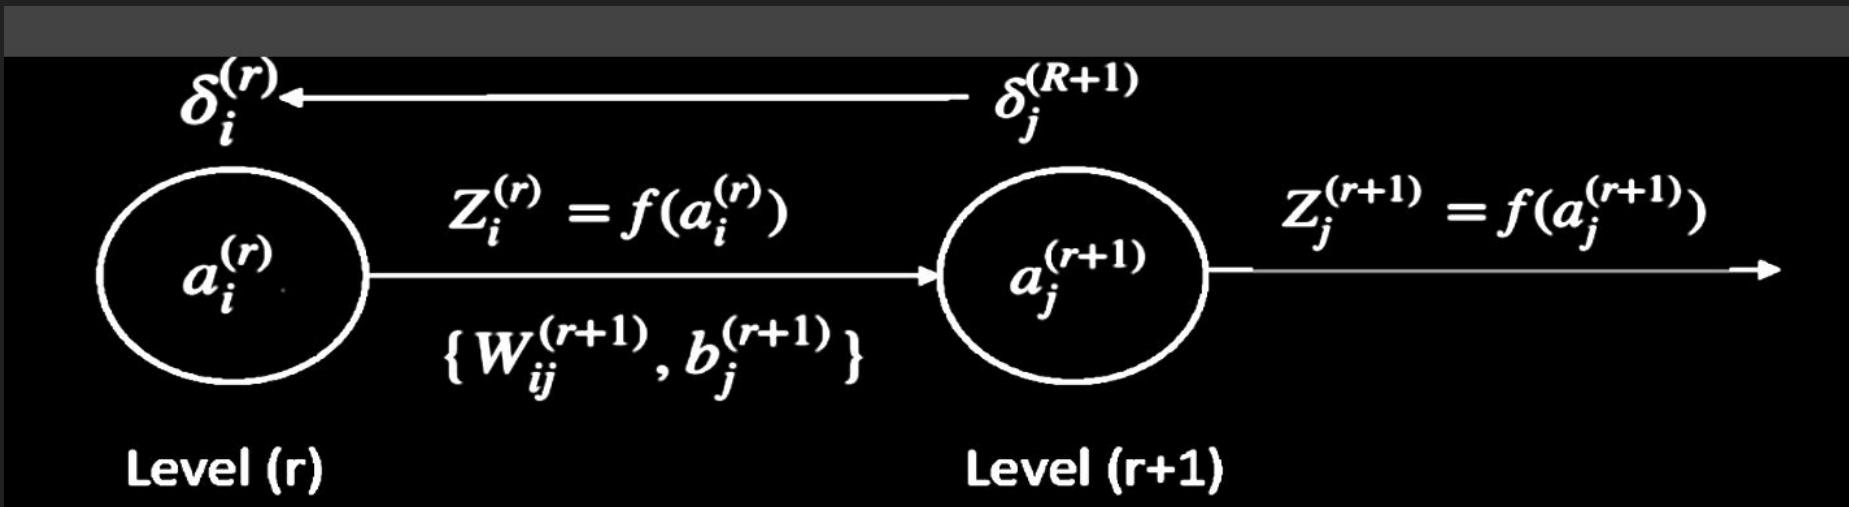
Final output = $h(a_j^{(R+1)})$

$$\delta_j^{(R+1)} = \frac{\partial L_m}{\partial a_j^{(R+1)}} = h'(a_j^{(R+1)})$$

Iterative Process

1. Start with an initial set of weights $\{w, b\}$.
2. Feedforward the initial data and weights into the DLN, and find the $\{a_i^{(r)}, Z_i^{(r)}\}, \forall r, i$.
3. Then, using backpropagation, compute all $\delta_i^{(r)}, \forall r, i$.
4. Use these $\delta_i^{(r)}$ values to get all the new gradients.
5. Apply gradient descent to get new weights.
6. Keep iterating steps 2-5, until the chosen number of epochs is completed.

Feedforward and Backprop



Recap

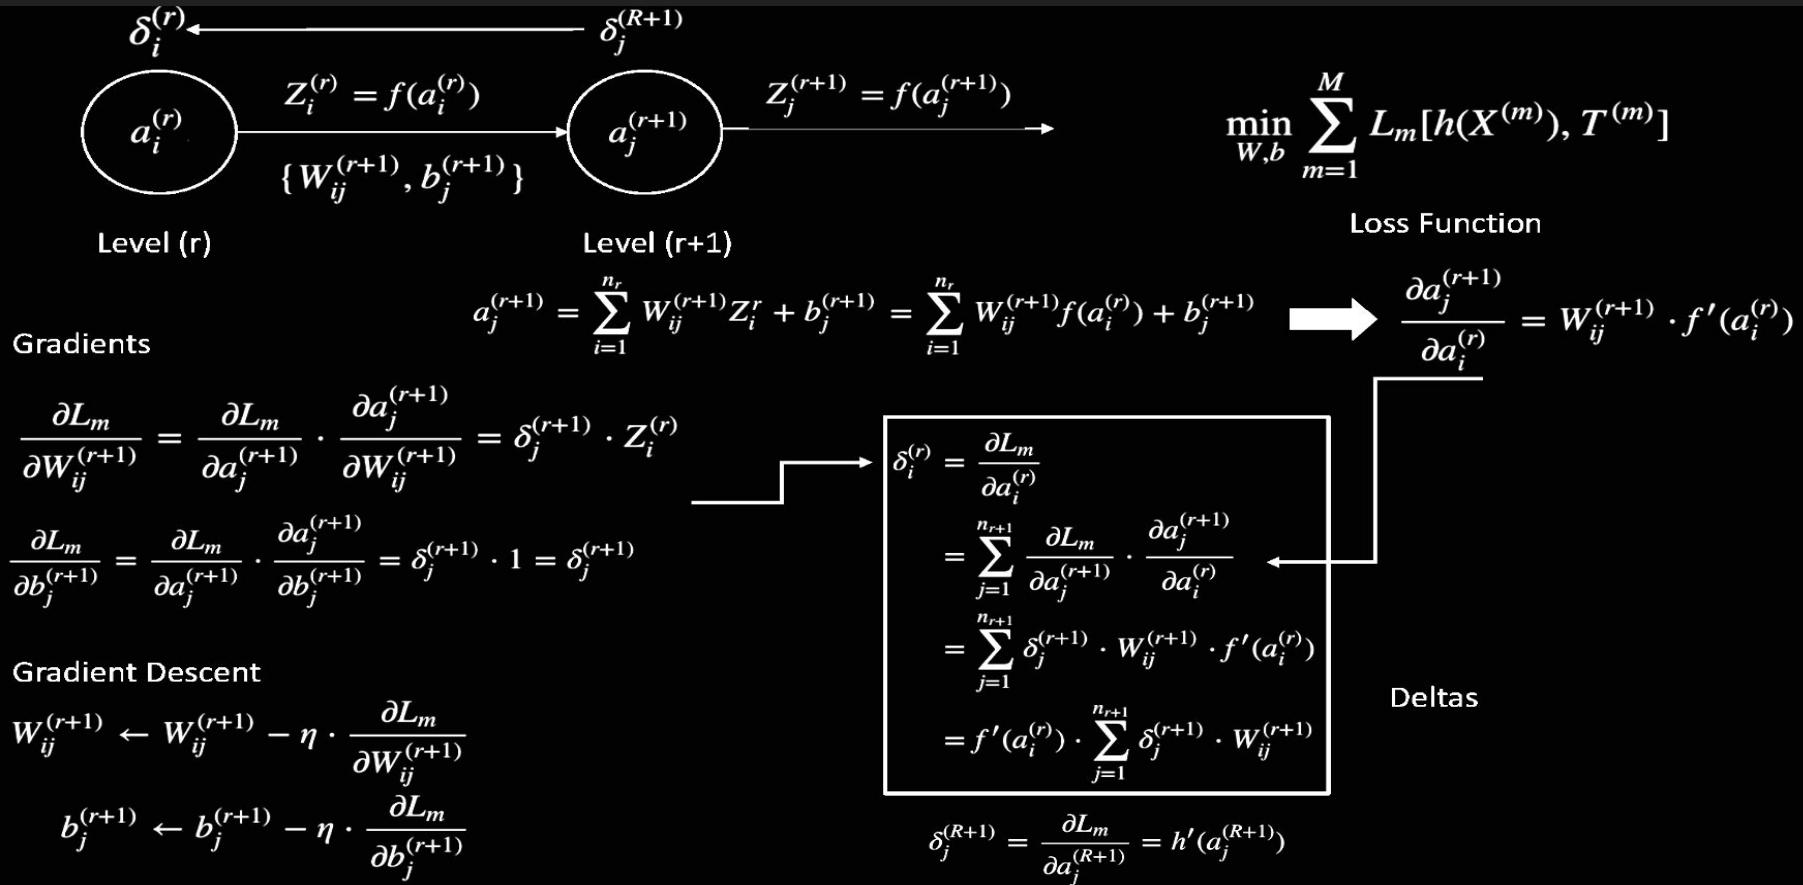
$$\frac{\partial L_m}{\partial W_{ij}^{(r+1)}} = \frac{\partial L_m}{\partial a_j^{(r+1)}} \cdot \frac{\partial a_j^{(r+1)}}{\partial W_{ij}^{(r+1)}} = \delta_j^{(r+1)} \cdot Z_i^{(r)}$$

$$\delta_j^{(r+1)} = \frac{\partial L_m}{\partial a_j^{(r+1)}}$$

$$= f'(a_i^{(r)}) \cdot \sum_{j=1}^{n_{r+1}} \delta_j^{(r+1)} \cdot W_{ij}^{(r+1)}$$

$$\frac{\partial L_m}{\partial b_j^{(r+1)}} = \frac{\partial L_m}{\partial a_j^{(r+1)}} \cdot \frac{\partial a_j^{(r+1)}}{\partial b_j^{(r+1)}} = \delta_j^{(r+1)} \cdot 1 = \delta_j^{(r+1)}$$

The Magic of Backpropagation



Genesis: Breast Cancer Detection

- We use a data set of breast cancer cell measurements and assess whether we can predict whether a tumor is malignant or benign.
- The data used comprises nine cell measurements such as cell thickness, dimension, etc., and from these measurements we wish to attain a high level of accuracy in classification of cells.
- A system such as this can potentially be the first line of diagnosis and may be able to improve the doctor's diagnosis, in a Bayesian sense, as there are now two readings of the cell, by the deep learning net and the doctor.
- The data set is the Wisconsin breast cancer data.

TensorFlow

TensorFlow™

Install

Develop

API r1.0

Deploy

Extend

Resources

Versions



Search

Github

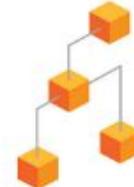
An open-source software library
for Machine Intelligence

GET STARTED



TensorFlow 1.0 has arrived!

We're excited to announce the release of TensorFlow 1.0! Check out the migration guide to upgrade your code with ease.



Dynamic graphs in TensorFlow

We've open-sourced TensorFlow Fold to make it easier than ever to work with input data with varying shapes and sizes.



The 2017 TensorFlow Dev Summit

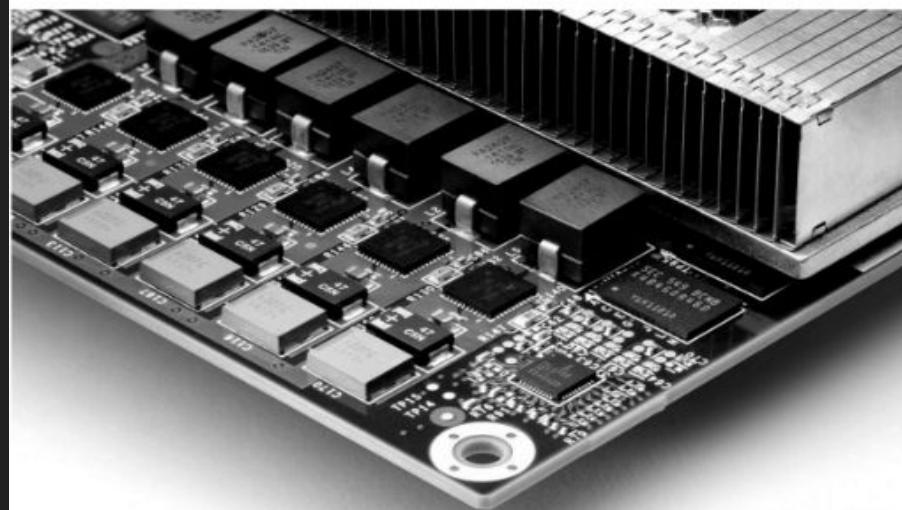
Thousands of people from the TensorFlow community participated in the first flagship event. Watch the keynote and talks.

Special Purpose Chips

Google's Tensor Processing Unit explained: this is what the future of computing looks like

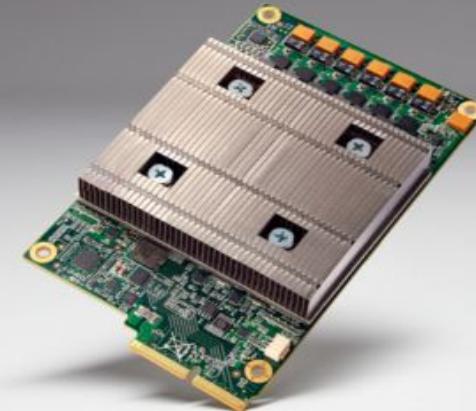
By Joe Osborne August 22, 2016 Processors

Specialized silicon will lead the way forward



Google's Big Chip Unveil For Machine Learning: Tensor Processing Unit With 10x Better Efficiency (Updated)

by Lucian Armasu May 19, 2016 at 9:10 AM - Source: Google Cloud Platform Blog



TensorFlow Playground

Epoch
000,000Learning rate
0.03Activation
TanhRegularization
NoneRegularization rate
0Problem type
Classification

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%

Noise: 0

Batch size: 10

REGENERATE

FEATURES

Which properties do you want to feed in?



+ - 2 HIDDEN LAYERS

+ -

4 neurons

+ -

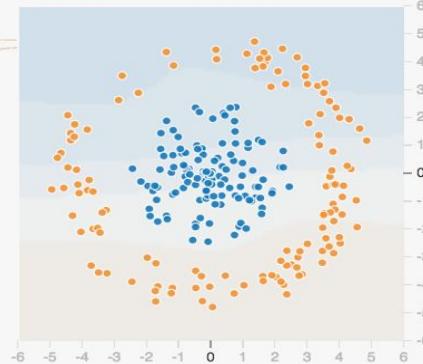
2 neurons

This is the output from one neuron. Hover to see it larger.

The outputs are mixed with varying weights, shown by the thickness of the lines.

OUTPUT

Test loss 0.500
Training loss 0.500



Colors shows data, neuron and weight values.



Show test data

Discretize output

<http://playground.tensorflow.org/>

Genesis: Breast Cancer Detection

- We use a data set of breast cancer cell measurements and assess whether we can predict whether a tumor is malignant or benign.
- The data used comprises nine cell measurements such as cell thickness, dimension, etc., and from these measurements we wish to attain a high level of accuracy in classification of cells.
- A system such as this can potentially be the first line of diagnosis and may be able to improve the doctor's diagnosis, in a Bayesian sense, as there are now two readings of the cell, by the deep learning net and the doctor.
- The data set is the Wisconsin breast cancer data.

Cancer Data

```
In [2]: %pylab inline  
import pandas as pd  
  
Populating the interactive namespace from numpy and matplotlib
```

```
In [4]: %% Read in the data set  
data = pd.read_csv("BreastCancer.csv")  
data.head()
```

```
Out[4]:
```

	Id	Cl.thickness	Cell.size	Cell.shape	Marg.adhesion	Epith.c.size	Bare.nuclei	Bl.cromatin
0	1000025	5	1	1	1	2	1	3
1	1002945	5	4	4	5	7	10	3
2	1015425	3	1	1	1	2	2	3
3	1016277	6	8	8	1	3	4	3
4	1017023	4	1	1	3	2	1	3

```
x = data.loc[:, 'Cl.thickness':'Mitoses']  
print(x.head())  
y = data.loc[:, 'Class']  
print(y.head())
```

	Cl.thickness	Cell.size	Cell.shape	Marg.adhesion	Epith.c.size
0	5	1	1	1	1
1	5	4	4	4	5
2	3	1	1	1	1
3	6	8	8	8	1
4	4	1	1	1	3

	Bare.nuclei	Bl.cromatin	Normal.nucleoli	Mitoses
0	1	3	1	1
1	10	3	2	1
2	2	3	1	1
3	4	3	7	1
4	1	3	1	1


```
0    benign  
1    benign  
2    benign  
3    benign  
4    benign  
Name: Class, dtype: object
```

<https://futurism.com/machine-learning-identifies-breast-lesions-likely-to-become-cancer/>

TensorFlow and One-Hot Encoding

```
## Convert the class variable into binary numeric
ynum = zeros((len(x),1))
for j in arange(len(y)):
    if y[j]=="malignant":
        ynum[j]=1
ynum[:10]
array([[ 0.],
       [ 0.],
       [ 0.],
       [ 0.],
       [ 0.],
       [ 1.],
       [ 0.],
       [ 0.],
       [ 0.],
       [ 0.]])
```

```
## Define the neural net and compile it
from keras.models import Sequential
from keras.layers import Dense, Activation

model = Sequential()
model.add(Dense(32, activation='relu', input_dim=9))
model.add(Dense(32, activation='relu'))
model.add(Dense(32, activation='relu'))
```

```
## Accuracy
yhat = model.predict_classes(x, batch_size=32)
acc = sum(yhat==ynum)
print("Accuracy = ", acc/len(ynum))

## Confusion matrix
from sklearn.metrics import confusion_matrix
confusion_matrix(yhat, ynum)
```

Using TensorFlow back

```
[[ 1.  0.]
 [ 1.  0.]
 [ 1.  0.]
 [ 1.  0.]
 [ 1.  0.]
 [ 0.  1.]
 [ 1.  0.]
 [ 1.  0.]
 [ 1.  0.]
 [ 1.  0.]]
(683, 9)
(683, 2)
(683, 1)
```

```
32/683 [>.....] - ETA: 0sAccuracy =  0.978038067
35
array([[431,   2],
       [ 13, 237]])
```

Canonical Example: Digit Recognition

THE MNIST DATABASE

of handwritten digits

Yann LeCun, Courant Institute, NYU

Corinna Cortes, Google Labs, New York

Christopher J.C. Burges, Microsoft Research, Redmond

The MNIST database of handwritten digits, available from this page, has a training set of 60,000 examples, and a test set of 10,000 examples. It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image.



C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	C11	C12	C13	C14	C15	C16	C17	C18	C19	C20	C21	C22	C23	C24	C25	C26	C27	C28
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
11	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
12	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
13	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
14	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
15	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
16	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
17	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
18	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
19	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
20	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
21	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
22	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
23	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
24	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
25	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

Showing 1 to 25 of 60,000 entries

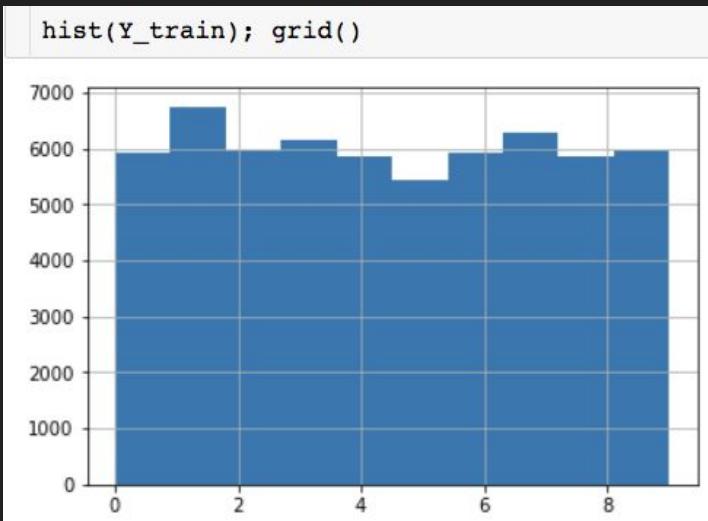
- Futile to run a multinomial regression on 784 variables
- Extensible to many finance problems

<https://www.cs.toronto.edu/~kriz/cifar.html>

MNIST Data

```
## Read in the data set
train = pd.read_csv("train.csv", header=None)
test = pd.read_csv("test.csv", header=None)
print(shape(train))
print(shape(test))

(60000, 785)
(10000, 785)
```



```
## Reformat the data
X_train = train.as_matrix()[:,0:784]
Y_train = train.as_matrix()[:,784:785]
print(shape(X_train))
print(shape(Y_train))
X_test = test.as_matrix()[:,0:784]
Y_test = test.as_matrix()[:,784:785]
print(shape(X_test))
print(shape(Y_test))
y.labels = utils.to_categorical(Y_train, num_classes=10)
print(shape(y.labels))
print(y.labels[1:5,:])
print(Y_train[1:5])

(60000, 784)
(60000, 1)
(10000, 784)
(10000, 1)
(60000, 10)
[[ 0.  0.  0.  1.  0.  0.  0.  0.  0.  0.]
 [ 1.  0.  0.  0.  0.  0.  0.  0.  0.  0.]
 [ 1.  0.  0.  0.  0.  0.  0.  0.  0.  0.]
 [ 0.  0.  1.  0.  0.  0.  0.  0.  0.  0.]]
[[3]
 [0]
 [0]
 [2]]
```

TF Model

```
## Define the neural net and compile it
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout
from keras.optimizers import SGD

data_dim = shape(X_train) #> ## Fit/train the model (x,y need to be matrices)
model.fit(X_train, y.labels, epochs=5, batch_size=32,verbose=2)

model = Sequential()
model.add(Dense(100, activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(10, activation='softmax'))

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

Accuracy

```
## In Sample
yhat = model.predict_classes(X_train, batch_size=32)

## Confusion matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(yhat,Y_train)
print(" ")
print(cm)

##
acc = sum(diag(cm))/len(Y_train)
print("Accuracy = ",acc)
```

```
59392/60000 [=====>.] - ETA: 0:00
[[5749    0   45   16   11   39   29   23   20   42]
 [  0  6533   23   12   15    4   11   24   61    6]
 [ 14   47 5568  167   19   27   14   59   55    5]
 [  8   22   48 5505   2  116    0   17   72   71]
 [ 16    4   45    4 5522   21   20   48   26  203]
 [ 26   23   11  185    5 5004   71    4   84   40]
 [ 42    8   87   16   70   84 5747    1   39    4]
 [  3   19   41   70    4   16    0 5982   11   76]
 [ 62   76   83  130   25   77   26    7 5435   89]
 [  3   10    7   26  169   33    0 100   48 5413]]
Accuracy =  0.940966666667
```

```
## Out of Sample
yhat = model.predict_classes(X_test, batch_size=32)

## Confusion matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(yhat,Y_test)
print(" ")
print(cm)

##
acc = sum(diag(cm))/len(Y_test)
print("Accuracy = ",acc)
```

```
9728/10000 [=====>.] - ETA: 0:00
[[ 959    0   11    1    2   12   16    2    7   10]
 [  0 1115    0    1    3    0    2   14    6    4]
 [  3    2   964   30    7    6    9   16   16    1]
 [  1    4   922    0   43    1    3   18    9]
 [  1    0    9    0 923    5    9    5   12   41]
 [  6    0    1   23    0 772   10    0   19    4]
 [  6    3   10    0   11   15 908    0   10    1]
 [  2    2   11   14    1    7    0 961    7   15]
 [  2    9   13   14    3   27    3    0 871    9]
 [  0    0    1    5   32    5    0   27    8 915]]
Accuracy =  0.931
```

Amazon Rekognition

<https://console.aws.amazon.com/rekognition/home?region=us-east-1#/label-detection>

The screenshot shows the Amazon Rekognition service page within the AWS Management Console. The top navigation bar includes links for Services, Resource Groups, Notifications (bell icon), User (Sanjiv Das), Region (N. Virginia), and Support. The main content area features a dark blue background with a network graph pattern. At the top center is the title "Amazon Rekognition" in large white font, followed by the subtitle "Deep learning-based image recognition service" and "Search, verify, and organize millions of images". Below the subtitle are two buttons: "Try Demo" and "Download SDKs". To the left, a sidebar menu lists various services and resources under "Amazon Rekognition": Metrics, Demos (Object and scene detection, Image moderation, Facial analysis, Celebrity recognition, Face comparison), Additional Resources (Getting started guide, Download SDKs, Developer resources), Pricing, FAQ, and Forum. The main content area also contains three sections with icons and descriptions: "Get Context and Information from Images" (stacked squares icon), "Scalable Deep Learning-based Image Analysis" (circuit board icon), and "Integrated with AWS Services" (puzzle pieces icon).

Amazon Rekognition

Deep learning-based image recognition service
Search, verify, and organize millions of images

Try Demo

Download SDKs

Get Context and Information from Images

Amazon Rekognition enables your applications to gain a high-level understanding of the content of images. With Rekognition, your applications can detect objects, scenes, and faces in images. With this rich context and information, you can enable visual search, discovery, and authentication in

Scalable Deep Learning-based Image Analysis

Amazon Rekognition is designed to use deep learning technology to analyze billions of images daily. The APIs find and compare faces, detect thousands of objects. We continue to add new objects and make improvements to facial analysis so you can focus on building applications.

Integrated with AWS Services

Amazon Rekognition is designed to work seamlessly with other AWS services. Rekognition integrates directly with Amazon S3 and AWS Lambda so you can build scalable, affordable, and reliable image analysis applications. You can start analyzing images stored in Amazon S3 without moving any data. Support for Identity and Access Management (IAM)

Learning the Black-Scholes Equation (Culkin & Das, 2017)

THE JOURNAL OF FINANCE • VOL. XLIX, NO. 3 • JULY 1994

A Nonparametric Approach to Pricing and Hedging Derivative Securities Via Learning Networks

JAMES M. HUTCHINSON, ANDREW W. LO, and
TOMASO POGGIO*

ABSTRACT

We propose a nonparametric method for estimating the pricing formula of a derivative asset using learning networks. Although not a substitute for the more traditional arbitrage-based pricing formulas, network-pricing formulas may be more accurate and computationally more efficient alternatives when the underlying asset's price dynamics are unknown, or when the pricing equation associated with the no-arbitrage condition cannot be solved analytically. To assess the potential value of network pricing formulas, we simulate Black-Scholes option prices and show that learning networks can recover the Black-Scholes formula from a two-year training set of daily options prices, and that the resulting network formula can be used successfully to both price and delta-hedge options out-of-sample. For comparison, we estimate models using four popular methods: ordinary least squares, radial basis function networks, multilayer perceptron networks, and projection pursuit. To illustrate the practical relevance of our network pricing approach, we apply it to the pricing and delta-hedging of S&P 500 futures options from 1987 to 1991.

```
from scipy.stats import norm
def BSM(S,K,T,Sig,rf,dv,cp): #cp = {+1.0 (calls), -1.0 (puts)}
    d1 = (math.log(S/K)+(rf-dv+0.5*sig**2)*T)/(sig*math.sqrt(T))
    d2 = d1 - sig*math.sqrt(T)
    return cp*S*math.exp(-dv*T)*norm.cdf(d1*cp) - cp*K*math.exp(-rf*T)*norm.cdf(d2*cp)

df = pd.read_csv('/Users/srdas/GoogleDrive/Papers/DeepLearning/DLinFinance/SP500Options.csv')
```

Normalizing spot and call prices

C is homogeneous degree one, so

$$aC(S, K) = C(aS, aK)$$

This means we can normalize spot and call prices and remove a variable by dividing by K .

$$\frac{C(S, K)}{K} = C(S/K, 1)$$

Data, libraries, activation functions

```
n = 300000
n_train = (int)(0.8 * n)
train = df[0:n_train]
x_train = train[['Stock Price', 'Maturity', 'Dividends', 'Volatility', 'Risk-Free Rate']]
y_train = train['Call Price'].values
test = df[n_train+1:]
x_test = test[['Stock Price', 'Maturity', 'Dividends', 'Volatility', 'Risk-Free Rate']]
y_test = test['Call Price'].values

#Import libraries
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, LeakyReLU
from keras import backend

def custom_activation(x):
    return backend.exp(x)
```

Fit the Model

```
nodes = 120
model = Sequential()

model.add(Dense(nodes, input_dim=X_train.shape[1]))
model.add(LeakyReLU())
model.add(Dropout(0.25))

model.add(Dense(nodes, activation='elu'))
model.add(Dropout(0.25))

model.add(Dense(nodes, activation='relu'))
model.add(Dropout(0.25))

model.add(Dense(nodes, activation='elu'))
model.add(Dropout(0.25))

model.add(Dense(1))
model.add(Activation(custom_activation))

model.compile(loss='mse',optimizer='rmsprop')

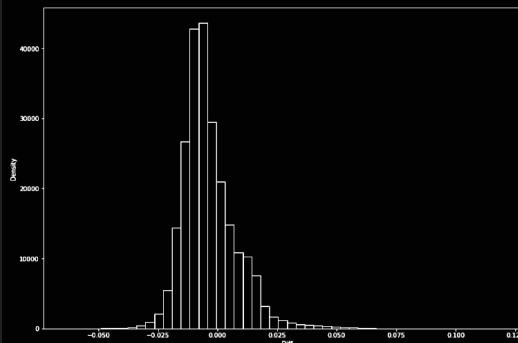
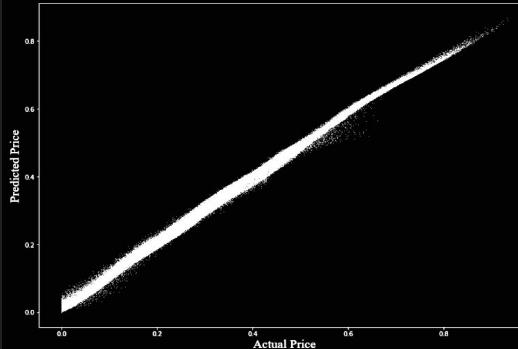
model.fit(X_train, y_train, batch_size=64, epochs=10, validation_split=0.1)

Train on 216000 samples, validate on 24000 samples
Epoch 1/10
22s - loss: 0.0051 - val_loss: 0.0013
Epoch 2/10
20s - loss: 0.0013 - val_loss: 1.3056e-04
Epoch 3/10
21s - loss: 0.0010 - val_loss: 7.0662e-04
Epoch 4/10
20s - loss: 8.4915e-04 - val_loss: 2.9489e-04
Epoch 5/10
21s - loss: 7.4240e-04 - val_loss: 6.5513e-04
Epoch 6/10
21s - loss: 6.7416e-04 - val_loss: 4.9628e-04
Epoch 7/10
21s - loss: 6.3666e-04 - val_loss: 3.7541e-04
Epoch 8/10
21s - loss: 5.9715e-04 - val_loss: 1.2603e-04
Epoch 9/10
26s - loss: 6.0532e-04 - val_loss: 1.3601e-04
Epoch 10/10
29s - loss: 6.1739e-04 - val_loss: 1.3957e-04
<keras.callbacks.History at 0x10ba4cef0>
```

In-Sample

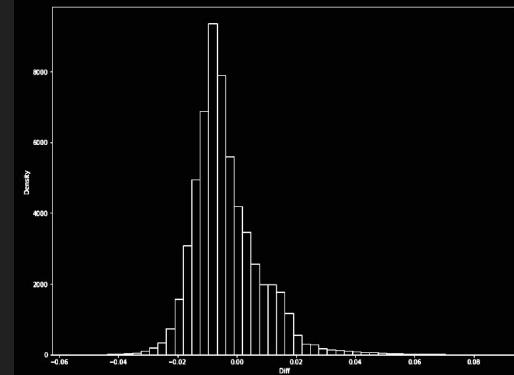
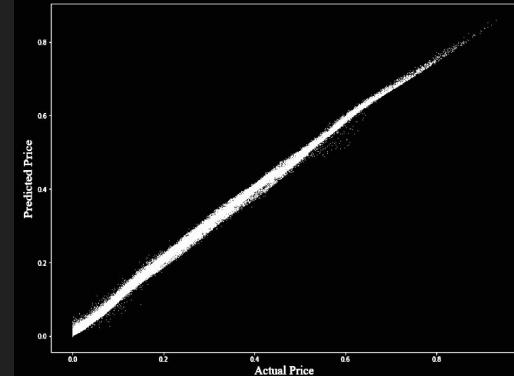
```
y_train_hat = model.predict(X_train)
#reduce dim (240000,1) -> (240000,) to match y_train's dim
y_train_hat = squeeze(y_train_hat)
CheckAccuracy(y_train, y_train_hat)
```

Mean Squared Error: 0.00014108053002
Root Mean Squared Error: 0.0118777325278
Mean Absolute Error: 0.00954217479875
Mean Percent Error: 0.0444015623019

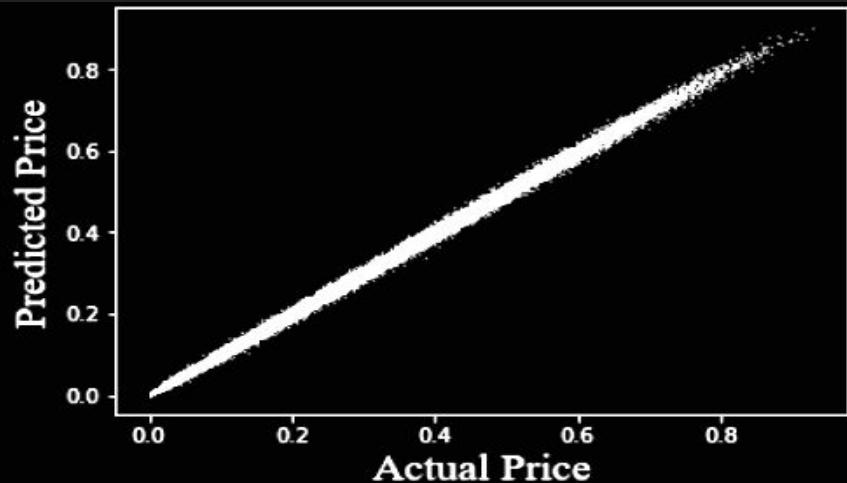
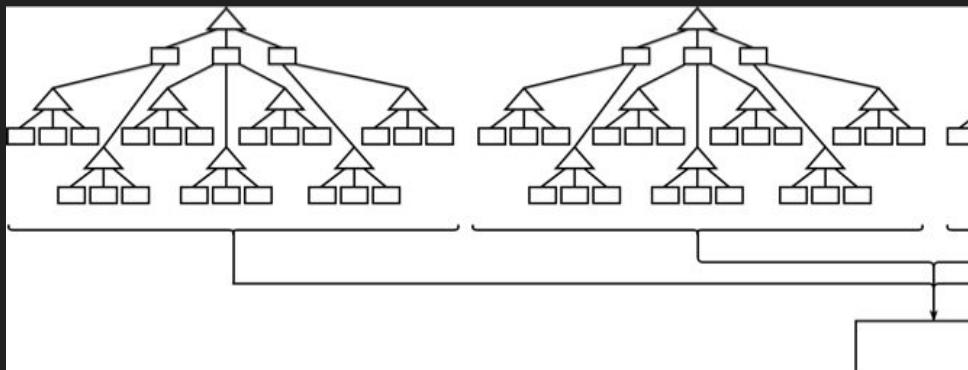


Out-of-Sample

Mean Squared Error: 0.00014210254092
Root Mean Squared Error: 0.0119240743423
Mean Absolute Error: 0.00958741885268
Mean Percent Error: 0.0446387882149



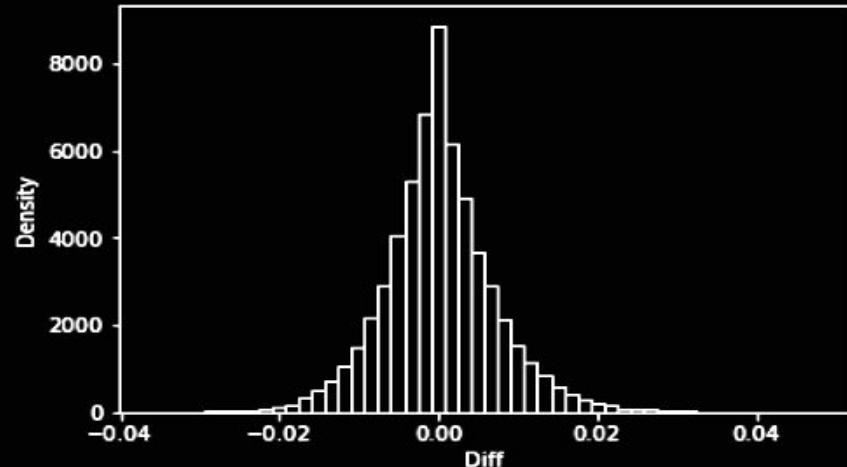
Random Forest



```
from sklearn.ensemble import RandomForestRegressor  
  
forest = RandomForestRegressor()  
forest = forest.fit(X_train, y_train)  
y_test_hat = forest.predict(X_test)
```

```
stats = CheckAccuracy(y_test, y_test_hat)
```

Mean Squared Error: 4.97733713185e-05
Root Mean Squared Error: 0.00705502454415
Mean Absolute Error: 0.00516501148658
Mean Percent Error: 0.0264110854593



Are Markets Still Efficient?



A new kind of hedge fund built by a network of data scientists.

[JOIN NUMERAI](#) [SIGN IN](#)

The stock market is inefficient with respect to new developments in machine learning – only a fraction of the world's data scientists have access to its data. Numerai is changing this.

Predicting the Direction of the S&P500

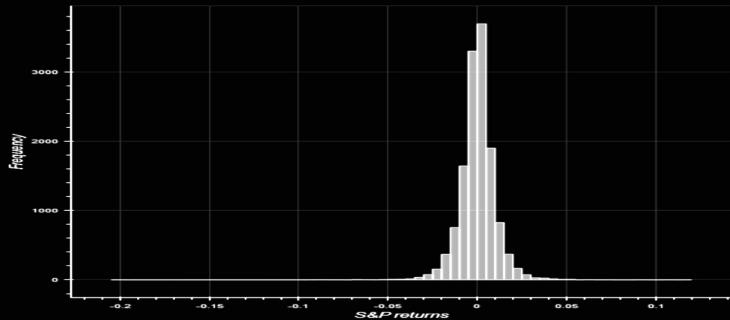
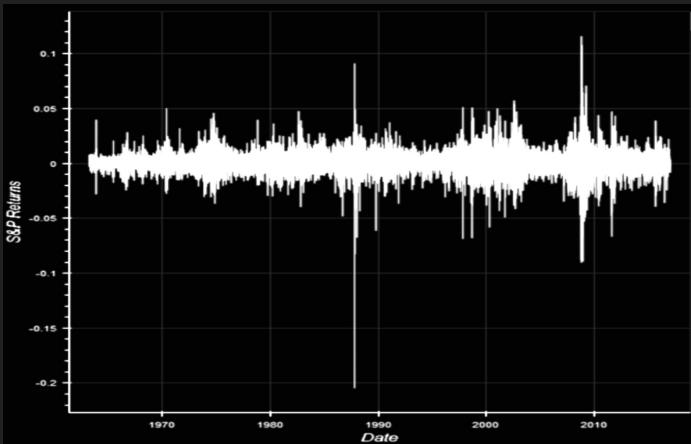


Figure 2: The distribution of daily S&P 500 index returns from 1963-2016. The mean return is 0.00031 and the standard deviation is 0.01. The skewness and kurtosis are -0.62 and 20.68, respectively.



Culkin, Das, Mokashi (2017)

Data, lookback, lookforward

```
n = 30          #Lookback period  
n fwd = 30      #Look ahead forecast period  
train_size = 10000    #Training size
```

```
m = shape(df)[0]  
print(m) #Number of rows
```

```
df2 = df[n:m][df.columns[[0,1,21]]]  
print(shape(df2))  
df2.head()
```

```
13532  
(13502, 3)
```

	date	SPX	Sign
30	1963-05-14	-0.003831	0.0
31	1963-05-15	0.003133	1.0
32	1963-05-16	-0.002556	0.0
33	1963-05-17	0.000569	1.0
34	1963-05-20	-0.004695	0.0

Organize data and fit the model

```
## Fit the model to the first 10,000 rows of data
X_train = df2[x:x+train_size].as_matrix()
X_train = X_train[:,3:len(df2.columns)]
print(shape(X_train))
Y_train = df2[x:x+train_size].as_matrix()
Y_train = Y_train[:,2:3].astype(int32)
print(shape(Y_train))

data_dim = shape(X_train)[1]
print(data_dim)

X_test = df2[x+train_size:(x+train_size)+test_size].as_matrix()
X_test = X_test[:,3:len(df2.columns)]
print(shape(X_test))
Y_test = df2[x+train_size:(x+train_size)+test_size].as_matrix()
Y_test = Y_test[:,2:3].astype(int32)
print(shape(Y_test))

(10000, 570)
(10000, 1)
570
(30, 570)
(30, 1)
```

```
model = Sequential()
n_units = 64 #200

model.add(Dense(n_units, input_dim=data_dim))
model.add(LeakyReLU())
model.add(Dropout(0.25))
```

Model Accuracy

$$\text{parameters} = 570 \times 65 + 64 \times 65 + 64 \times 65 + 65 = 45435$$

```
model.evaluate(X_train, Y_train2, verbose=0)
```

```
[0.68329057350158695, 0.5565499999999999]
```

```
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout
from keras.layers.advanced_activations import LeakyReLU
from keras.utils import to_categorical

Y_train2 = to_categorical(Y_train, 2)
```

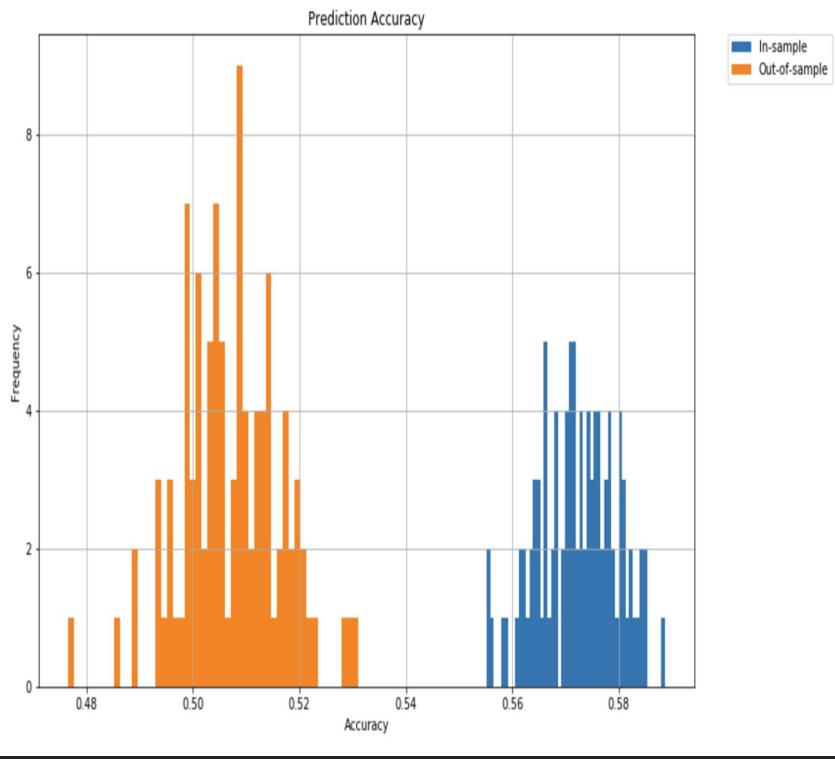
```
loss='binary_crossentropy',
metrics=['accuracy'])

bsize = 32
model.fit(X_train, Y_train2, batch_size=bsize, epochs=25, validation_
```

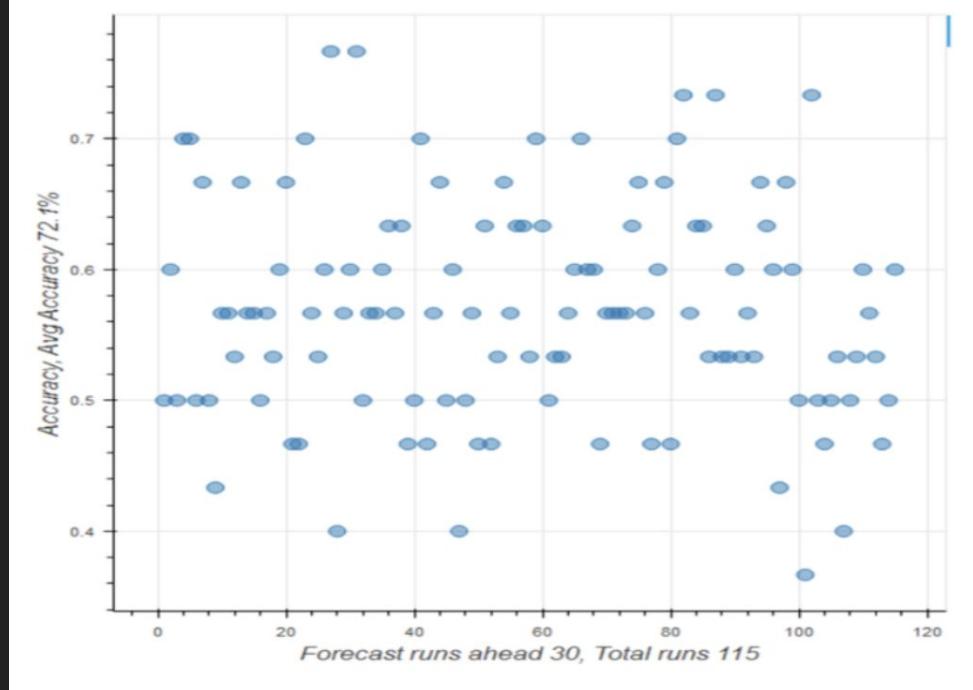
Train on 9000 samples, validate on 1000 samples

Epoch 1/25
2s - loss: 0.6939 - acc: 0.5124 - val_loss: 0.6924 - val_acc: 0.5180
Epoch 2/25
2s - loss: 0.6935 - acc: 0.5171 - val_loss: 0.6899 - val_acc: 0.5430
Epoch 3/25
2s - loss: 0.6936 - acc: 0.5156 - val_loss: 0.6916 - val_acc: 0.5160

Predictions



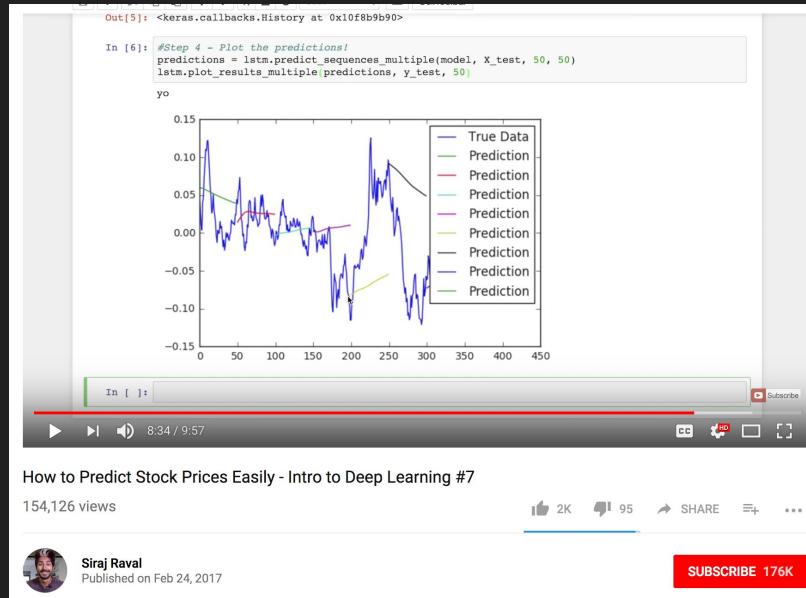
TensorFlow



H20.deeplearning : lookback 30 days, forward predict 30 days, total experiments 115.

- More than 50% right 72.1% of the time
- Overall accuracy 54.4%

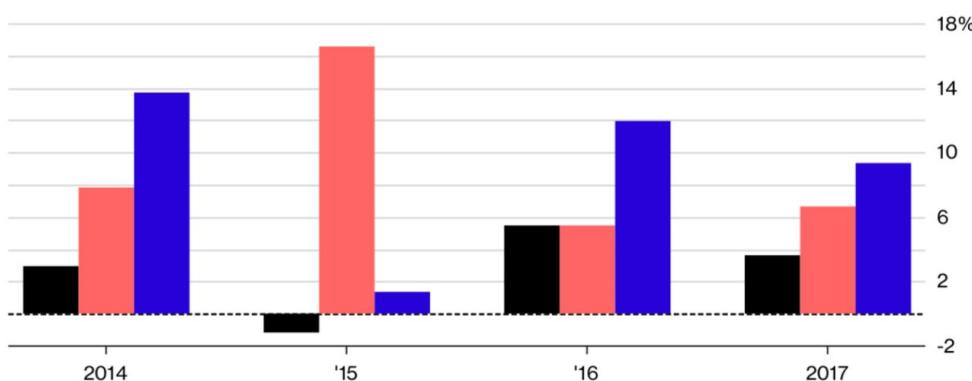
Everyone's trying to deep learn market prediction



Machine Learning's Gains

Like hedge funds, AI strategies have struggled to beat the stock market

■ HFRI Fund Weighted Composite Index ■ Eurekahedge AI Hedge Fund Index ■ S&P 500



2017 returns YTD through June, S&P 500 Index returns are with dividend reinvested
Source: Eurekahedge, Hedge Fund Research, Inc., Bloomberg

Bloomberg

<https://www.youtube.com/watch?v=ftMq5ps503w>

Two Curses of Predictive Analytics

Non-stationarity

- Strong: Joint distribution of all variables (Y, X) remains the same over time.
- Weak: only mean and autocorrelation need to be same over time. If we are predicting the first moment, then this works.

Randomness

- If noise swamps the signal, then we get poor predictions.

In short, stock-picking is hard:

<https://www.wsj.com/articles/the-future-is-bumpy-high-tech-hedge-fund-hits-limits-of-robot-stock-picking-1513007557>

Or is it?

<https://www.bloomberg.com/news/articles/2017-12-11/hedge-fund-founder-and-amazon-researcher-build-a-better-ai-stock-picker>

Estimating the effects of deep learning

- (Roy) Amara's Law: "We tend to overestimate the effect of a technology in the short run and underestimate the effect in the long run."
- Arthur C. Clarke's Three Laws:
 - a. When a distinguished but elderly scientist states that something is possible, he is almost certainly right. When he states that something is impossible, he is very probably wrong.
 - b. The only way of discovering the limits of the possible is to venture a little way past them into the impossible.
 - c. Any sufficiently advanced technology is indistinguishable from magic.

Results of Market Prediction Analysis

Add more slides here on using Python for market prediction analysis.

AI as a Service

- Commoditized Platforms: Amazon's AWS Machine Learning VMs.
- Commoditized Services: e.g., Image Rekognition. Document analysis.
- Commoditized Models: Model zoos. NLP offerings.
- Encrypted Crowdsourcing of AI. e.g., Numerai. Managing data and privacy.
- Federated Machine Learning. Distributed Learning and data privacy.

Federated Machine Learning

INNOVATION

How Google wants to crowdsource machine learning with smartphones and Federated Learning

Google's new approach to machine learning relies on data stored on mobile devices to train models, without the need to centrally store the data.

By Conner Forrest | April 7, 2017, 11:22 AM PST



SAP® Basis Survival Guide

Learn How To Manage SAP Like A BOSS With This Guide. Get It Now!

WHITE PAPERS, WEBCASTS, AND DOWNLOADS

White Papers // From Adobe Systems

Mobile Learning Leading The Charge

The learning expectations of today's workers have changed. Traditional classroom and eLearning courses have less relevance in today's corporate, fast-paced environment. Our research shows that online learning is as effective as – and in some cases more effective than – classroom learning. Technology enables fast, efficient learning and allows for in-depth analysis and problem-based learning. Digital disruption affects the tools, content, and delivery methods that make up your...

DOWNLOAD NOW

White Papers // From GoTransverse

A leader in Recurring Customer And Billing Management

DOWNLOAD NOW

Training // From IBM

Social media understood: What IBM Watson Analytics for Social Media can do for you

GIVE IT A TRY!

White Papers // From IBM

Other applications of deep learning in finance

- Forecasting the VIX curve.
- Predicting Credit Card Default, see Khandani, Lee, Lo (2016).
- Portfolio optimization.
- Text analytics.
- Text generation (e.g., Narrative Science).

What happens to financial sector employment?



Elon Musk 
@elonmusk

 Follow

Replying to @elonmusk

China, Russia, soon all countries w strong computer science.

Competition for AI superiority at national level most likely cause of WW3 imo.

2:33 AM - Sep 4, 2017

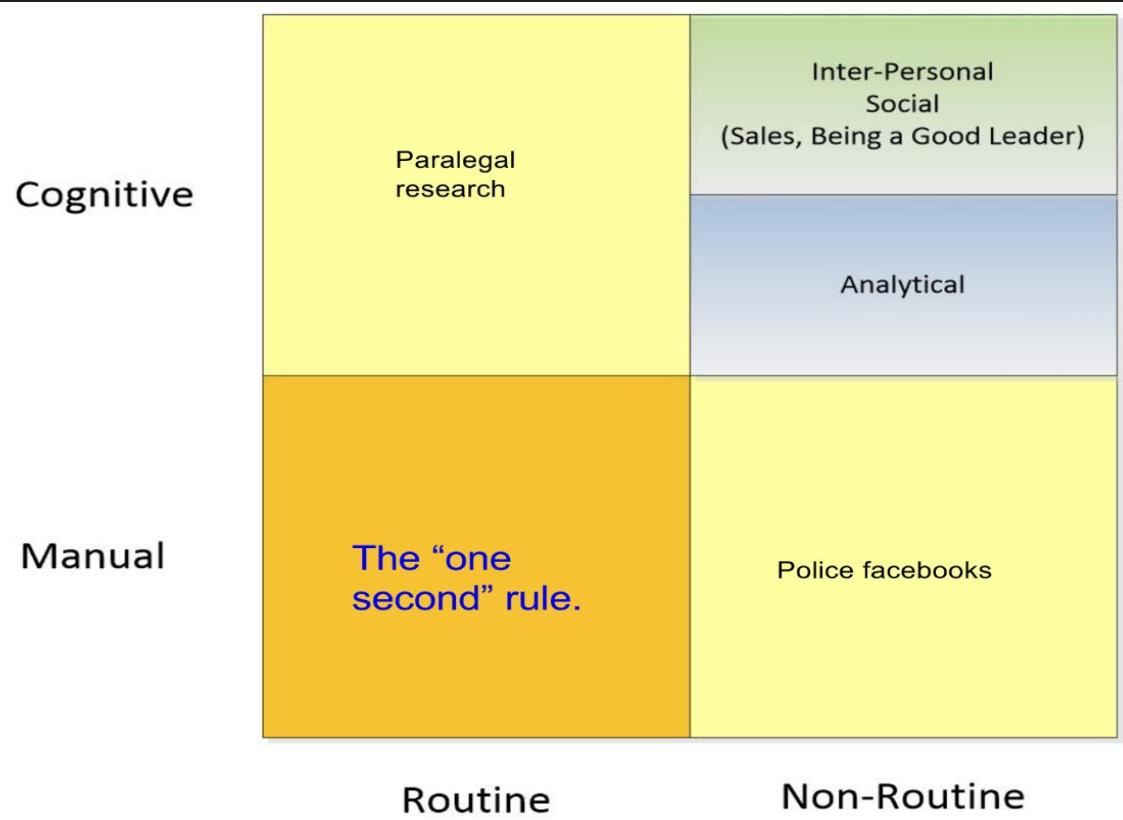
 3,540

 19,169

 46,908



The Atomic Level of Work



What tasks get automated first?

The End !!

Readings on AI and Deep Learning

1. <http://srdas.github.io/DLBook/>
2. <https://medium.com/@katherinebailey/hashtag-artificial-intelligence-47ff35e6a9cc>
3. <https://www.linkedin.com/pulse/ai-deep-learning-explained-simply-fabio-ciucci>
4. <https://hackernoon.com/why-ai-is-now-on-the-menu-at-dinner-even-with-my-non-tech-friends-44c666348de4>
5. <https://futurism.com/kurzweil-claims-that-the-singularity-will-happen-by-2045/>
6. <https://www.quantamagazine.org/artificial-intelligence-learns-to-learn-entirely-on-its-own-20171018/>
7. https://www.technologyreview.com/s/603984/googles-ai-explosion-in-one-chart/?utm_campaign=add_this&utm_source=twitter&utm_medium=post
8. <http://www.zerohedge.com/news/2017-02-13/goldman-had-600-cash-equity-traders-2000-it-now-has-2>
9. <https://futurism.com/an-ai-completed-360000-hours-of-finance-work-in-just-seconds/>
10. <https://www.theguardian.com/technology/2016/dec/22/bridgewater-associates-ai-artificial-intelligence-management>
11. https://www.mizuhobank.com/mizuho_fintech/news/pepper/index.html
12. <http://www.businessinsider.com/royal-bank-of-scotland-launches-ai-chatbot-luvo-using-ibm-watson-2016-9?r=UK&IR=T&IR=T>
13. <https://www.the-digital-insurer.com/dia/xtra-by-axa-ai-driven-personal-wellness-coaching-app/>
14. <http://www.nanalyze.com/2017/04/ai-fintech-startups-loans-new-credit/>
15. <https://www.americanbanker.com/news/this-is-how-financial-services-chatbots-are-going-to-evolve>
16. <http://fortune.com/2017/03/30/blackrock-robots-layoffs-artificial-intelligence-ai-hedge-fund/>
17. https://en.wikipedia.org/wiki/Sentient_Technologies
18. <https://medium.com/numeraian/an-ai-hedge-fund-goes-live-on-ethereum-a80470c6b681>
19. <https://venturebeat.com/2017/05/06/ai-powered-trading-raises-new-questions/>
20. <http://www.wired.co.uk/article/how-ai-is-transforming-the-future-of-fintech>
21. <https://medium.com/numeraian/encrypted-data-for-efficient-markets-ffffbe9743ba8>
22. <https://www.bloomberg.com/news/features/2017-09-27/the-massive-hedge-fund-betting-on-ai>
23. https://www.wired.com/2017/02/ai-threat-isnt-skynet-end-middle-class/?mbid=social_twitter_onsiteshare
24. <https://www.wired.com/2008/06/pb-theory/>
25. <https://www.gartner.com/smarterwithgartner/gartner-predicts-our-digital-future/>